

Advancing Legal Reasoning in *Algerian Law*

RAG · KNOWLEDGE GRAPHS · ARGUMENT MINING → *Citation-faithful legal answers*Ibrahim El Khalil Attia · Maab Chaoui | Supervisor **Dr. Fouad Dahak** Co-supervisor **Prof. Ivan Cantador**

01 The Problem

WHY WHY LEGAL QA IS HARD

1.1 Fluent is not lawful

A legal answer cannot be judged by fluency or plausibility. To be valid it must be **grounded in identifiable sources**, cite the **exact governing articles**, stay inside the **Algerian jurisdiction**, and read each provision **as amended** at the relevant date — and where the law is silent, the only correct answer is a **refusal**. A fluent language model guarantees none of this.

SOURCED · ARTICLE-CITED · JURISDICTION-BOUND · AMENDMENT-AWARE · OR *REFUSED*

1.2 What legal questions demand

يَبَيِّنُ الأساس القانوني لمبدأ خضوع القاضي للنص التشريعي، وحدود سلطته عند غياب النص...

— demands **exact articles** — *Civil Code, arts. 1–2*

كيف تغيّرت شروط الخلع في قانون الأسرة بعد تعديل 2005؟

— demands *the law as amended* — *Ord. 05-02 → Family Code, art. 54*

هل يحق لصاحب العمل في الجزائر فصل بحرية تامة دون إبداء أسباب (at-will)؟

— a **foreign-doctrine trap** — *Algerian law is silent; the system must abstain**verbatim from AlgerianLegalBench (244 questions)*

1.3 Where today's LLMs fail




Fig. 1. Hallucinated-citation rate: the share of cited articles that *do not exist*, per direct-LLM baseline (244 questions). Scale does not help — the 70B model fabricates 4 of 5 citations. Gemini 2.5 reaches 0.008 only by declining to cite at all, and six of seven baselines import foreign law on >12 of 40 trap questions. **AKN-RLM: 0.000.**

1.4 ...and Algerian law raises the bar

- **Arabic-first.** 232 of 244 benchmark questions are in Arabic
- **Layered in time.** 1963–2025; amendments rewrite articles in place
- **Deeply structured.** code → chapter → section → article
- **Low-resource.** barely covered by existing legal-AI corpora

1.5 Objectives

- ✓ **Structure the law**
171 docs → *Akoma Ntoso XML* + legal KG
- ✓ **Benchmark it**
AlgerianLegalBench: 244 expert questions
- ✓ **Reason on evidence**
typed retrieval + bounded recursion
- ✓ **Verify or abstain**
triple faithfulness gate

1.6 Significance & impact

1. **First citation-faithful QA system for Algerian law**
— every claim resolves to a registered article, or the system abstains.
2. **Reusable infrastructure for Algerian legal AI**
 - **AKN corpus + knowledge graph** — 171 structured documents, 1963–2025, article-level registry
 - **AlgerianLegalBench** — 244 expert-annotated questions, $\kappa = 0.829$
3. **A recipe for low-resource legal AI** — evidence

02 The Solution

HOW STRUCTURED CORPUS → BENCHMARK → AKN-RLM

I Structured legal corpus

the law, made machine-readable

What is Akoma Ntoso? The OASIS LegalDocML standard — an XML vocabulary for legal texts in which every law, article, amendment and cross-reference becomes an **addressable node** under a canonical FRBR URI: the article registry that every citation must resolve against.

WHAT WE BUILT

171 source files 189 MB 45 laws & codes → AKN XML 46 RDF graphs → one legal KG 1963–2025 coverage, amendments in place


```
<!-- Civil Code — Ordinance 75-58, as amended -->
<act name="code_civil">
  <FRBRWork value="/akn/dz/act/1975-09-26/75-58"/>
  <FRBRExpression value=".../ara@2005-06-20"/>
  <lifecycle>
    <eventRef date="2005-06-20" type="amendment"
      source="#law-05-10"/>
  </lifecycle>
  <article eId="art_124">
    <num>124</num>
    <content>كل فعل أيًا كان يرتكبه الشخص بخطئه</content>
    <ref href="#art_127"/>
  </article>
</act>
```

Validated in **multiple manual passes** by two law students and a supervising senior legal expert — the structure is derived from the sources, never invented.


II AlgerianLegalBench

the evaluation layer

244 expert-annotated questions across 23 legal categories — 177 answerable, 67 expecting **abstention** · 232 Arabic / 12 French · each carries expected articles, a reference answer and a **four-step reasoning chain** · double-annotated, $\kappa = 0.829$.



III AKN-RLM

the answer pipeline — Fig. 3.2

03 Results

PROOF FROZEN 244-QUESTION RUN

0.305 Citation F1
article-level · locked


0.000 HCR · JIR

0.703 Abstention F1

1.64×
*vs best direct LLM*1.74×
*vs minimal RAG*2.9×
vs determ. retrieval

+0.199 over the strongest same-corpus baseline, $p_{Holm} < 0.0001$ ·
document-level F1 0.546


rlm_dispatched_full_phase_e_final · 2026-05-14 · depth ≤ 3 · NLI ≥ 0.55 · one retry → abstain



Each phase is an ablation in disguise — **every layer pays for itself**, though not always in F1.


IV Citation F1 vs baselines

HIGHER = BETTER



V Jurisdictional infection (JIR)

LOWER = SAFER



Only AKN-RLM holds **0.000** hallucinated-citation & foreign-law rates. Reliability tracks **legal structure**, not **scale** — the largest baselines hallucinate most.

VI Reading the results

INTERPRETATION

- The gain is architectural**
identical corpus, registry & indices as every tier-2 baseline — the layers, not the scale, earn the F1
- Strictness has a price**
the gate over-rejects layman queries; wider evidence pools dilute citation discipline
- Refusal works**
38 / 40 infection traps correctly refused — every baseline invents an answer

LIMITS

- Retrieval ceiling: $R@10 = 0.258$
- French under-served
- Abstention precision 0.654
- Metrics are

NEXT

- Corpus-tuned embedder
- French ↔ Arabic retrieval
- Graph-retrieval redesign



READ THE THESIS

Built on a structured Algerian legal corpus — Akoma Ntoso XML + Knowledge Graph



OUR CVS